

## Product Management Endpoints for a Retail/E-commerce Application


<http://assignment-react-app.s3-website-ap-southeast-1.amazonaws.com/>

Github source code: [nhanne/ProductService](https://github.com/nhanne/ProductService)

## Approach

Understand the requirements, design the domain models, then build with Clean Architecture + DDD.

## Database



- SQL (PostgreSQL) because product data needs ACID and consistent relationships.
- Flexible attributes stored in JSONB or Variant table → can add new product attributes anytime without schema changes.
- Split read/write DB later for performance. On read side, disable EF change tracking to speed up queries.

## Tech Stack

- .NET 8, Entity Framework Core (Code First)
- MediatR (CQRS), FluentValidation (input validation)
- Serilog (logs), Autofac (DI)

- API Design: Standard RESTful CRUD, Swagger (API docs), Use DTOs for requests/responses.
- Validate: input with FluentValidation, Always return global exception on error (400/404/500).

## Caching

I use Redis for caching, and basically there are two common strategies:

### 1. Lazy Loading (Cache-Aside)

- On read: check cache → if miss, load from DB → put result into cache.
- Cheap and simple, but may have stale data right after a write.

### 2. Write-Through

- Whenever the database is updated, I also update the cache immediately.
- Ensures the cache always has the latest data, but it increases write cost and latency.

My Approach: Hybrid

- Frequently updated or critical data → Write-Through (Keeps cache fresh and avoids stale reads).
- Less frequently updated data → Lazy Loading (Saves cost and still improves performance).

## Concurrency

There are three ways to handle concurrent updates:

### 1. Optimistic Locking (use RowVersion)

- No lock, just check version when saving.
- Fast and scalable, but conflicts require retry.

### 2. Pessimistic Locking (SELECT ... FOR UPDATE)

- Lock row until transaction finishes.
- No retry needed, but blocks other requests and reduces throughput.

### 3. Queue / FIFO Processing

- Put requests into a queue and process one by one in background.
- Handles huge traffic, avoids conflicts, but is eventually consistent.

## Query Performance

AsNoTracking(), indexes, and read-only repository to handle queries traffic.

## Limitation

Even though I don't have too much time to implement full caching and concurrency logic, I clearly understand the best practices and the trade-offs for each approach (optimistic/pessimistic locking, queueing, lazy-loading cache vs write-through, etc.).

## Future Imporve

### Unit Testing

Write automated unit tests to cover all business logic and edge cases.

### Image Handling

Move image storage to **Amazon S3** with signed URLs and proper lifecycle policies.

### Monitoring & Search


Implement full logging/monitoring with Serilog + Kibana.

Build a search engine for products using **Elasticsearch**.

### Deployment & Infrastructure

Currently: Client deployed to S3, API deployed as **Docker container on ECS**.

Future: Improve **networking, security groups, and standardize infrastructure** for better security and maintainability.



eraser

All diagram: <https://app.eraser.io/workspace/ZuxJp0sPbJ4zaCz4Pmvn?origin=share>